YAX Referee-Revision Diagnosis

Evidence-led framing decision, completion status, and unresolved claims

September 2026

Decision

YAX is a viable measurement-centered chapter. It is not a causal paper on AI displacement and should not be sold as general robustness of the early-career employment decline. The recommended title is the adopted title: *Constructed Exposure Measures and Statement-Specific Robustness: Evidence from Early-Career Employment*.

The revised contribution is stronger because it states exactly where conclusions change. Taxonomy repair mainly readmits occupations. Tail estimates are stable within two selected families but attenuate under two external concepts. Q1 is a consequential reference. Age and era definitions matter in specific directions. Continuous family asymmetry survives alpha deletion but becomes imprecise under other task-family deletions. Mobility ranking disagreement is common for small movements yet has little effect on aggregate upward/downward shares.

Evidence that forced the reframing

1. **Architecture scope.** The common-support six-measure result covers two dependent families. Webb AI estimates −0.0649 and OECD capability exposure −0.0110, both imprecise. The original all-negative statement cannot be generalized.
2. **Reference dependence.** Q5–Q1 is −0.1311, but Q5–Q2 is −0.0456 and Q5–Q4 −0.0340, with both intervals spanning zero. Favorable Q1 evolution is part of the economic contrast.
3. **Specificity.** Beta is more negative than education, cognitive content, and telework on one common sample, but a within-SOC2 permutation gives lower-tail probability .219. AI specificity is not established.
4. **Comparison ages and time.** Ages 22–25 versus 26–35 reproduce the pooled result; comparison with 51–65 is smaller. The 2025–26 coefficient is more negative than 2024, but not than 2023. The era-average pattern is nonmonotone.
5. **Family outcome sensitivity.** No-alpha F/G remains signed and precise, but G intervals cross zero after omitting beta or broad. The conclusion is construction-sensitive asymmetry, not a universal family ranking.
6. **Mobility scale.** Conflict is 53.28 percent at zero, 40.10 percent under movement-mass weighting, and substantial opposition is 14.00 percent at .5 SD. Aggregate moving-higher shares stay around

one half.

Evidence that remains useful

The mapping contribution is concrete: exact-code AIOE coverage of computer/math employment is 3.33 percent and the documented repair reaches 97.7 percent. The four-row coefficient decomposition shows that score correction on fixed support is small, while re-admission changes the estimand. The calendar repair restores March basic samples and leaves the coefficient similar. Minimum-size, broad-occupation, approximate leisure/hospitality, stable-taxonomy, and post-2020 sensitivities preserve a negative beta tail estimate under their stated scopes.

The design discipline remains an asset. The predeclared simultaneous all-six band fails even though an IUT gives $p = .045$. The paired beta–alpha interval is not reinterpreted as equivalence. Realized standard errors are 3.65 and 3.17 times planning values. Failed seasonality and full-refit bootstrap attempts remain visible. These facts make the paper more credible when framed as an audit rather than a protected headline.

What the chapter may say

It may say that occupational exposure constructions differ in coverage, tail membership, residual information, and mobility rankings; that the selected six implementations share negative common-support point estimates; that the primary beta tail contrast survives named data and sample checks; that outside architectures attenuate the CPS contrast; and that ranking disagreement depends strongly on movement scale.

It may not say that AI caused employment loss, that the result is robust to AI-exposure definition generally, that Q5 alone deteriorated, that a failed detection establishes equivalence, that rematching proves excess conflict negligible, or that CPS exposure variation explains the difference between U.S. payroll and Danish adoption-linked evidence.

Completion status

The baseline was reproduced; protected tags were verified; calendar data were repaired; all four decisive checks were run; inference, information, precision, mobility, taxonomy, and external-architecture audits were executed; result receipts and hashes were saved; figures and tables were regenerated; and the manuscript, appendix, response, and diagnosis were rendered as separate PDFs.

Two empirical diagnostics failed honestly. The occupation-specific age-season model did not converge with 5,148 added nuisance parameters. Literal residual-wild full refits had zero admissible groupedbinomial pseudo-outcomes in 199 attempts. Neither failure was replaced with an outcome-favorable procedure.

The complete original referee-report files remain unavailable. The response matrix covers every comment in the master prompt but cannot certify word-for-word reconciliation of unseen minor text.

Author actions before journal submission

Complete affiliation, email, acknowledgments, funding, disclosure, and any data-use language required by the target journal. Supply the two original reports for a final wording reconciliation. Decide whether to retain the ReStat target after reading the revised measurement-centered paper; no journal-specific novelty claim has been invented. Do not edit the empirical claims without regenerating the final consistency audit.